



GPT Q1

LX / Tech: Sprache und Technologie - Language and Technology

niemand sollte umsonst arbeiten

Date: 2026-05-20

class: [16830] Sprache und Technologie - Language and Technology:: tutor: frankowsky :: term: SS26

snc: 16175.1

snc: 16132. 1 f d : 8 wer es [pɛlɛχɪl] ausspricht, und nicht [pɛrɛχɪl]

TO GIT COMMIT

[and wwell if youre desperate about not working links in the downloaded pdf please consider as intended to switch to the online version of this document witch you will find linked near header sorry your inconvenience is my convenience...]

[1] "engine on..."

snc

- 16215.1.2

random chat...

Q: chatGPT and Schwarz (2026)

16214.gpt.q

Q

wie war nochmal die frage?:)

A

This is one of the central epistemic problems for future language models: recursive synthetic training.

A simplified loop looks like this:

1. humans write internet text
2. models train on human text
3. models generate new internet text
4. future models train on mixtures of human + model text
5. the proportion of synthetic text increases each generation

Eventually:

- the training corpus becomes partially self-referential.

That has deep consequences.

1. Collapse of provenance

Historically, web text had an implicit anchoring:

- journalism
- books
- institutional records
- human testimony
- scientific publication
- forum discussions tied to lived experience

Even when wrong, those originated from some contact with reality.

Synthetic text weakens that anchoring.

Future corpora may contain:

- paraphrases of paraphrases
- generated summaries of generated summaries
- hallucinations repeated across thousands of sites
- mutually reinforcing but source-free claims

The model then learns: not “what happened,” but “what text statistically survived circulation.” Truth becomes partially detached from origin.

2. Epistemic feedback loops

Suppose an early model hallucinates:

“Scientist X proposed theory Y in 1974.”

Blogs copy it. SEO pages reproduce it. AI summaries restate it. Aggregators absorb it.

A later model sees:

- many independent-looking confirmations.

Statistically, the claim now *looks true*.

This creates:

- synthetic consensus.

The dangerous part: the model’s confidence calibration improves *internally* while correspondence to reality worsens externally.

3. Model collapse

There is already literature on this (“model collapse”).

If you repeatedly train on synthetic outputs:

- rare information disappears
- variance shrinks
- distributions flatten
- edge cases vanish
- language becomes more generic and self-similar

Like making a photocopy of a photocopy repeatedly.

Over generations:

- entropy structure degrades,
- not increases.

The model converges toward:

- highly probable,
- averaged,
- stereotyped representations.

Unexpected truths disappear first.

4. Distortion of minority reality

Human reality is uneven and messy.

Synthetic corpora preferentially amplify:

- dominant narratives
- SEO-optimized formulations
- culturally central viewpoints
- statistically common phrasing

Marginal experiences become underrepresented.

Future models may therefore:

- speak more smoothly,
- while knowing less about actual diversity of lived reality.

5. Separation between coherence and truth

This is perhaps the deepest issue.

Language models optimize:

- predictive coherence, not:
- ontological truth.

As synthetic text increases, future models may become extraordinarily coherent while drifting away from empirical grounding.

You could get propositions that are:

- internally reinforced,
- stylistically authoritative,
- widely repeated, yet false.

The danger is not obvious nonsense. The danger is plausible epistemic simulation.

6. Emergence of “consensus reality generators”

If enough institutions use model-generated summaries:

- search engines,
- news digests,
- educational material,
- corporate reports,
- legal drafting,

then models stop merely describing consensus.

They begin manufacturing it.

That changes the role of AI from: epistemic mirror to epistemic infrastructure.

7. Possible countermeasures

Future systems will likely require stronger grounding mechanisms:

Provenance tracking

Cryptographic/source lineage:

- human-authored?
- model-generated?
- citation chain?

Trusted corpora

Heavily curated datasets:

- scientific archives
- verified records
- primary documents

Retrieval over generation

Dynamic lookup becomes more important than static memorization.

Synthetic data filtering

Detect and down-weight AI-generated text.

Reality anchoring

Connection to:

- measurements,
 - sensors,
 - experiments,
 - human verification.
-

8. Philosophical implication

There is a strange inversion here.

Human language evolved because organisms needed coordination about reality.

But sufficiently recursive synthetic language can become:

- self-sustaining without reality contact.

A future model might become excellent at predicting: what other models would say, rather than: what the world is like.

In that regime, “truth” risks being replaced by: stable linguistic equilibrium.

references

chatGPT, and St. Schwarz. 2026. “How Recent Chats and Discourse Model Future LLM Output.” In *Cite Yourself Dummy. A Transformer Conversation.*, 1–16214. Berlin. <https://stschwarz.userpage.fu-berlin.de/links?title=gpt-con&ref=cite>.