



GPT Q1

LX / Tech: Sprache und Technologie - Language and Technology

niemand sollte umsonst arbeiten

Date: 2026-05-20

class: [16830] Sprache und Technologie - Language and Technology:: **tutor:** frankowsky :: **term:** SS26

snc: 16175.1

snc: 16132. 1] ɔ : ʁ wer es [pɛlɛχɪl] ausspricht, und nicht [pɛrɛχɪl]

Table of contents

random chat...	2
16214.gpt.q	2
Q	2
A	2
1. Collapse of provenance	3
2. Epistemic feedback loops	3
3. Model collapse	4
4. Distortion of minority reality	4
5. Separation between coherence and truth	5
6. Emergence of “consensus reality generators”	5
7. Possible countermeasures	6
Provenance tracking	6
Trusted corpora	6
Retrieval over generation	6
Synthetic data filtering	6
Reality anchoring	7
8. Philosophical implication	7
references	7

TO GIT COMMIT

[1] "engine on..."

snc

- 16215.1.2

random chat...

Q: chatGPT and Schwarz (2026)

16214.gpt.q

Q

wie war nochmal die frage?:)

A

This is one of the central epistemic problems for future language models: recursive synthetic training.

A simplified loop looks like this:

1. humans write internet text
2. models train on human text
3. models generate new internet text
4. future models train on mixtures of human + model text
5. the proportion of synthetic text increases each generation

Eventually:

- the training corpus becomes partially self-referential.

That has deep consequences.

1. Collapse of provenance

Historically, web text had an implicit anchoring:

- journalism
- books
- institutional records
- human testimony
- scientific publication
- forum discussions tied to lived experience

Even when wrong, those originated from some contact with reality.

Synthetic text weakens that anchoring.

Future corpora may contain:

- paraphrases of paraphrases
- generated summaries of generated summaries
- hallucinations repeated across thousands of sites
- mutually reinforcing but source-free claims

The model then learns: not “what happened,” but “what text statistically survived circulation.”

Truth becomes partially detached from origin.

2. Epistemic feedback loops

Suppose an early model hallucinates:

“Scientist X proposed theory Y in 1974.”

Blogs copy it. SEO pages reproduce it. AI summaries restate it. Aggregators absorb it.

A later model sees:

- many independent-looking confirmations.

Statistically, the claim now *looks true*.

This creates:

- synthetic consensus.

The dangerous part: the model's confidence calibration improves *internally* while correspondence to reality worsens externally.

3. Model collapse

There is already literature on this ("model collapse").

If you repeatedly train on synthetic outputs:

- rare information disappears
- variance shrinks
- distributions flatten
- edge cases vanish
- language becomes more generic and self-similar

Like making a photocopy of a photocopy repeatedly.

Over generations:

- entropy structure degrades,
- not increases.

The model converges toward:

- highly probable,
- averaged,
- stereotyped representations.

Unexpected truths disappear first.

4. Distortion of minority reality

Human reality is uneven and messy.

Synthetic corpora preferentially amplify:

- dominant narratives
- SEO-optimized formulations
- culturally central viewpoints
- statistically common phrasing

Marginal experiences become underrepresented.

Future models may therefore:

- speak more smoothly,
 - while knowing less about actual diversity of lived reality.
-

5. Separation between coherence and truth

This is perhaps the deepest issue.

Language models optimize:

- predictive coherence, not:
- ontological truth.

As synthetic text increases, future models may become extraordinarily coherent while drifting away from empirical grounding.

You could get propositions that are:

- internally reinforced,
- stylistically authoritative,
- widely repeated, yet false.

The danger is not obvious nonsense. The danger is plausible epistemic simulation.

6. Emergence of “consensus reality generators”

If enough institutions use model-generated summaries:

- search engines,
- news digests,
- educational material,
- corporate reports,
- legal drafting,

then models stop merely describing consensus.

They begin manufacturing it.

That changes the role of AI from: epistemic mirror to epistemic infrastructure.

7. Possible countermeasures

Future systems will likely require stronger grounding mechanisms:

Provenance tracking

Cryptographic/source lineage:

- human-authored?
- model-generated?
- citation chain?

Trusted corpora

Heavily curated datasets:

- scientific archives
- verified records
- primary documents

Retrieval over generation

Dynamic lookup becomes more important than static memorization.

Synthetic data filtering

Detect and down-weight AI-generated text.

Reality anchoring

Connection to:

- measurements,
 - sensors,
 - experiments,
 - human verification.
-

8. Philosophical implication

There is a strange inversion here.

Human language evolved because organisms needed coordination about reality.

But sufficiently recursive synthetic language can become:

- self-sustaining without reality contact.

A future model might become excellent at predicting: what other models would say, rather than: what the world is like.

In that regime, “truth” risks being replaced by: stable linguistic equilibrium.

references

chatGPT, and St. Schwarz. 2026. “How Recent Chats and Discourse Model Future LLM Output.” In *Cite Yourself Dummy. A Transformer Conversation.*, 1–16214. Berlin. <https://stschwarz.userpage.fu-berlin.de/links?title=gpt-con&ref=cite>.